

# 1\_Data>Loading\_and\_EDA (1)

December 12, 2025

## 0.1 1. Data Loading and Exploratory Data Analysis

```
[3]: import boto3
import zipfile
import io
import pandas as pd
# import matplotlib.pyplot as plt
# import seaborn as sns
import numpy as np
```

### 0.1.1 1.1 Data Loading

We have stored the zip file containing the data in S3, we will load it and unzip it into the Jupyter Lab environment

```
[1]: !aws s3 cp "s3://tfm-annagm/ieee-fraud-detection.zip" ./
/bin/bash: switchml: line 1: syntax error: unexpected end of file
/bin/bash: error importing function definition for `switchml'
/bin/bash: module: line 1: syntax error: unexpected end of file
/bin/bash: error importing function definition for `module'
download: s3://tfm-annagm/ieee-fraud-detection.zip to ./ieee-fraud-detection.zip
```

```
[4]: zip_path = "./ieee-fraud-detection.zip"
extract_to = "./input_data"

with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_to)
```

### 0.1.2 1.2 Exploratory Data Analysis

We have loaded all the files and we see two types of .csv for training: \* `train_identity.csv` \* `train_transaction.csv`

We will perform the Exploratory Data Analysis to both

```
[2]: transaction_df = pd.read_csv("./input_data/train_transaction.csv")
transaction_df.head(10)
```

```
[2]:   TransactionID  isFraud  TransactionDT  TransactionAmt  ProductCD  card1 \
0      2987000       0        86400          68.5           W  13926
1      2987001       0        86401          29.0           W  2755
2      2987002       0        86469          59.0           W  4663
3      2987003       0        86499          50.0           W  18132
4      2987004       0        86506          50.0           H  4497
5      2987005       0        86510          49.0           W  5937
6      2987006       0        86522          159.0          W  12308
7      2987007       0        86529          422.5          W  12695
8      2987008       0        86535          15.0           H  2803
9      2987009       0        86536          117.0          W  17399

    card2  card3      card4  card5  ...  V330  V331  V332  V333  V334  V335 \
0    NaN  150.0  discover  142.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
1  404.0  150.0  mastercard  102.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
2  490.0  150.0      visa  166.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
3  567.0  150.0  mastercard  117.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
4  514.0  150.0  mastercard  102.0  ...  0.0  0.0  0.0  0.0  0.0  0.0
5  555.0  150.0      visa  226.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
6  360.0  150.0      visa  166.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
7  490.0  150.0      visa  226.0  ...  NaN  NaN  NaN  NaN  NaN  NaN
8 100.0  150.0      visa  226.0  ...  0.0  0.0  0.0  0.0  0.0  0.0
9 111.0  150.0  mastercard  224.0  ...  NaN  NaN  NaN  NaN  NaN  NaN

    V336  V337  V338  V339
0  NaN  NaN  NaN  NaN
1  NaN  NaN  NaN  NaN
2  NaN  NaN  NaN  NaN
3  NaN  NaN  NaN  NaN
4  0.0  0.0  0.0  0.0
5  NaN  NaN  NaN  NaN
6  NaN  NaN  NaN  NaN
7  NaN  NaN  NaN  NaN
8  0.0  0.0  0.0  0.0
9  NaN  NaN  NaN  NaN

[10 rows x 394 columns]
```

```
[3]: identity_df = pd.read_csv("./input_data/train_identity.csv")
identity_df.head(10)
```

```
[3]:   TransactionID  id_01      id_02  id_03  id_04  id_05  id_06  id_07  id_08 \
0      2987004     0.0    70787.0    NaN    NaN    NaN    NaN    NaN    NaN
1      2987008    -5.0    98945.0    NaN    NaN    0.0   -5.0    NaN    NaN
2      2987010    -5.0   191631.0    0.0    0.0    0.0    0.0    NaN    NaN
3      2987011    -5.0   221832.0    NaN    NaN    0.0   -6.0    NaN    NaN
4      2987016     0.0    7460.0     0.0    0.0    1.0    0.0    NaN    NaN
```

```

5      2987017   -5.0    61141.0     3.0    0.0    3.0    0.0    NaN    NaN
6      2987022  -15.0      NaN    NaN    NaN    NaN    NaN    NaN    NaN
7      2987038    0.0    31964.0     0.0    0.0    0.0   -10.0    NaN    NaN
8      2987040  -10.0   116098.0     0.0    0.0    0.0    0.0    NaN    NaN
9      2987048   -5.0   257037.0     NaN    NaN    0.0    0.0    NaN    NaN

      id_09 ...          id_31  id_32          id_33          id_34  id_35 \
0      NaN ...  samsung browser 6.2    32.0  2220x1080  match_status:2    T
1      NaN ...  mobile safari 11.0    32.0  1334x750  match_status:1    T
2      0.0 ...       chrome 62.0    NaN      NaN      NaN      NaN    F
3      NaN ...       chrome 62.0    NaN      NaN      NaN      NaN    F
4      0.0 ...       chrome 62.0    24.0  1280x800  match_status:2    T
5      3.0 ...       chrome 62.0    24.0  1366x768  match_status:2    T
6      NaN ...           NaN    NaN      NaN      NaN      NaN    NaN
7      0.0 ...       chrome 62.0    32.0  1920x1080  match_status:2    T
8      0.0 ...       chrome 62.0    NaN      NaN      NaN      NaN    F
9      NaN ...       chrome 62.0    NaN      NaN      NaN      NaN    F

      id_36  id_37  id_38 DeviceType          DeviceInfo
0      F      T      T   mobile  SAMSUNG SM-G892A Build/NRD90M
1      F      F      T   mobile           iOS Device
2      F      T      T  desktop        Windows
3      F      T      T  desktop        NaN
4      F      T      T  desktop        MacOS
5      F      T      T  desktop        Windows
6      NaN    NaN    NaN      NaN        NaN
7      F      T      T   mobile        NaN
8      F      T      T  desktop        Windows
9      F      T      T  desktop        Windows

[10 rows x 41 columns]

```

```
[4]: train_df = pd.merge(transaction_df, identity_df, on='TransactionID', how='left')
```

We display basic information on the dataset merged below. And we see how not all transactions have identity information

```
[22]: print("Shape of transaction data:", transaction_df.shape)
print("Shape of identity data:", identity_df.shape)
print("Shape after merge:", train_df.shape)
train_df.head()
```

```
Shape of transaction data: (590540, 394)
Shape of identity data: (144233, 41)
Shape after merge: (590540, 434)
```

```
[22]: TransactionID  isFraud  TransactionDT  TransactionAmt  ProductCD  card1 \
0            2987000         0            86400          68.5          W    13926
```

```

1      2987001      0      86401      29.0      W    2755
2      2987002      0      86469      59.0      W    4663
3      2987003      0      86499      50.0      W   18132
4      2987004      0      86506      50.0      H   4497

      card2  card3      card4  card5 ...      id_31  id_32 \
0      NaN  150.0  discover  142.0 ...      NaN      NaN
1  404.0  150.0  mastercard  102.0 ...      NaN      NaN
2  490.0  150.0       visa  166.0 ...      NaN      NaN
3  567.0  150.0  mastercard  117.0 ...      NaN      NaN
4  514.0  150.0  mastercard  102.0 ...  samsung browser  6.2   32.0

      id_33      id_34  id_35  id_36  id_37  id_38  DeviceType \
0      NaN      NaN      NaN      NaN      NaN      NaN      NaN
1      NaN      NaN      NaN      NaN      NaN      NaN      NaN
2      NaN      NaN      NaN      NaN      NaN      NaN      NaN
3      NaN      NaN      NaN      NaN      NaN      NaN      NaN
4  2220x1080  match_status:2      T      F      T      T      mobile

      DeviceInfo
0      NaN
1      NaN
2      NaN
3      NaN
4  SAMSUNG SM-G892A Build/NRD90M

[5 rows x 434 columns]

```

```
[28]: missing = train_df.isnull().sum() / len(train_df) * 100
missing = missing[missing > 0].sort_values(ascending=False)
print("Columns with missing values:")
print(missing.head(100))
```

```
Columns with missing values:
id_24    99.196159
id_25    99.130965
id_07    99.127070
id_08    99.127070
id_21    99.126393
...
V236    77.913435
V237    77.913435
V240    77.913435
V249    77.913435
V252    77.913435
Length: 100, dtype: float64
```

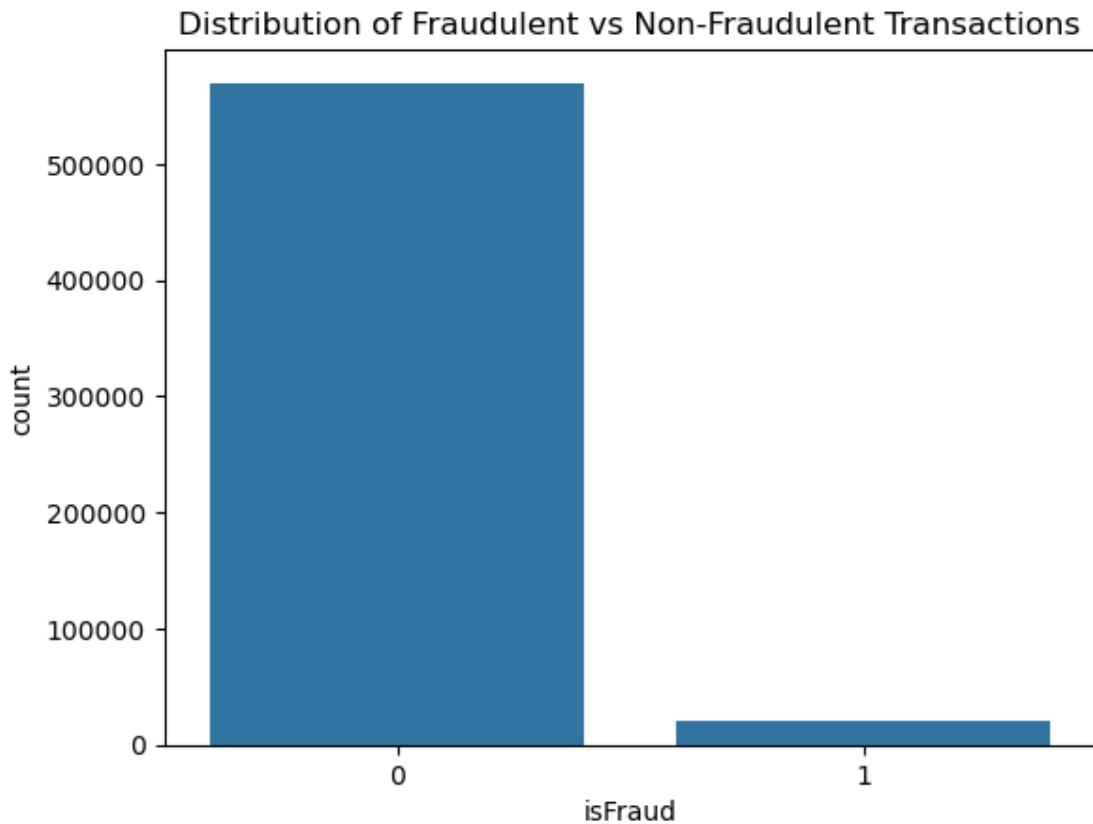
Many columns have a large number of nulls, so we will need to take that into account. If it

represents the data we will see in real transactions, it might be interesting to keep it as the model will need to learn from those nulls as well.

```
[33]: fraud_ratio = train_df['isFraud'].value_counts(normalize=True) * 100
print(fraud_ratio)

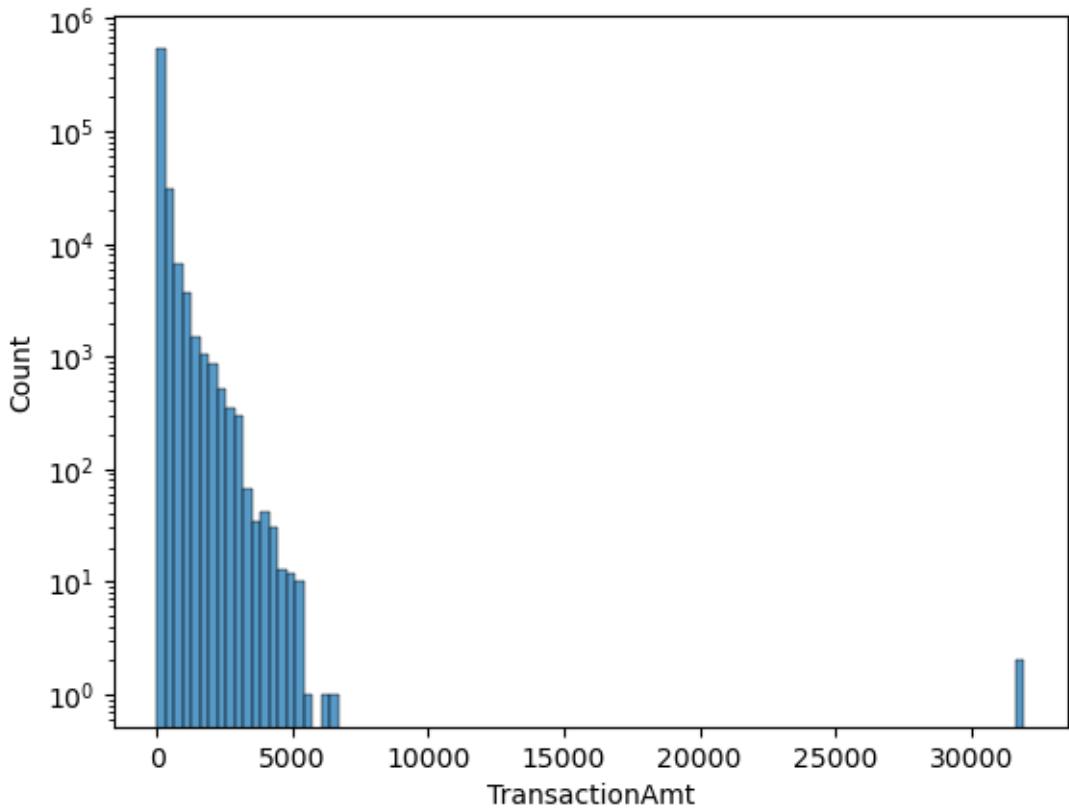
sns.countplot(data=train_df, x='isFraud')
plt.title('Distribution of Fraudulent vs Non-Fraudulent Transactions')
plt.show()
```

```
isFraud
0    96.500999
1     3.499001
Name: proportion, dtype: float64
```



We observe that the proportion of non-fraudulent cases is only 3.5%, which represents the reality of cases that have this imbalance in the datasets.

```
[36]: sns.histplot(train_df['TransactionAmt'], bins=100, log_scale=(False, True))
plt.yscale('log')
```



We can see that the transaction amounts are highly skewed, meaning that the majority of transactions are really small. This leads to a long right tail in the histogram. We observe it in logarithmic scale to be able to observe the large transactions as well.

```
[37]: plt.figure(figsize=(10,5))
sns.boxplot(data=train_df, x='isFraud', y='TransactionAmt')
plt.yscale('log')
plt.title('Transaction Amount vs Fraud')
plt.show()
```



If we observe the Transaction Amount for fraudulent vs non-fraudulent cases we can see that fraudulent cases shows a slightly higher median and a wider interquartile range (IQR). This suggests that while most frauds happen at small amounts (likely to avoid detection), some frauds involve large transactions, creating the long upper tail.

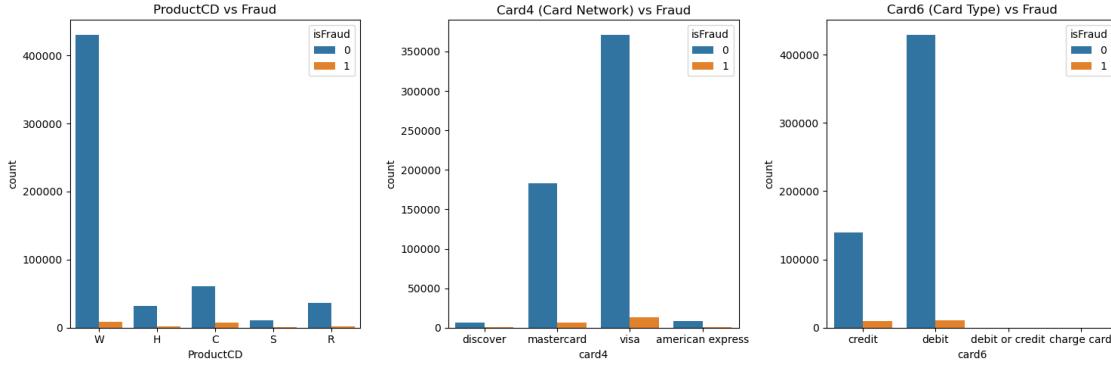
```
[39]: fig, axes = plt.subplots(1, 3, figsize=(15,5))

sns.countplot(data=train_df, x='ProductCD', hue='isFraud', ax=axes[0])
axes[0].set_title('ProductCD vs Fraud')

sns.countplot(data=train_df, x='card4', hue='isFraud', ax=axes[1])
axes[1].set_title('Card4 (Card Network) vs Fraud')

sns.countplot(data=train_df, x='card6', hue='isFraud', ax=axes[2])
axes[2].set_title('Card6 (Card Type) vs Fraud')

plt.tight_layout()
plt.show()
```

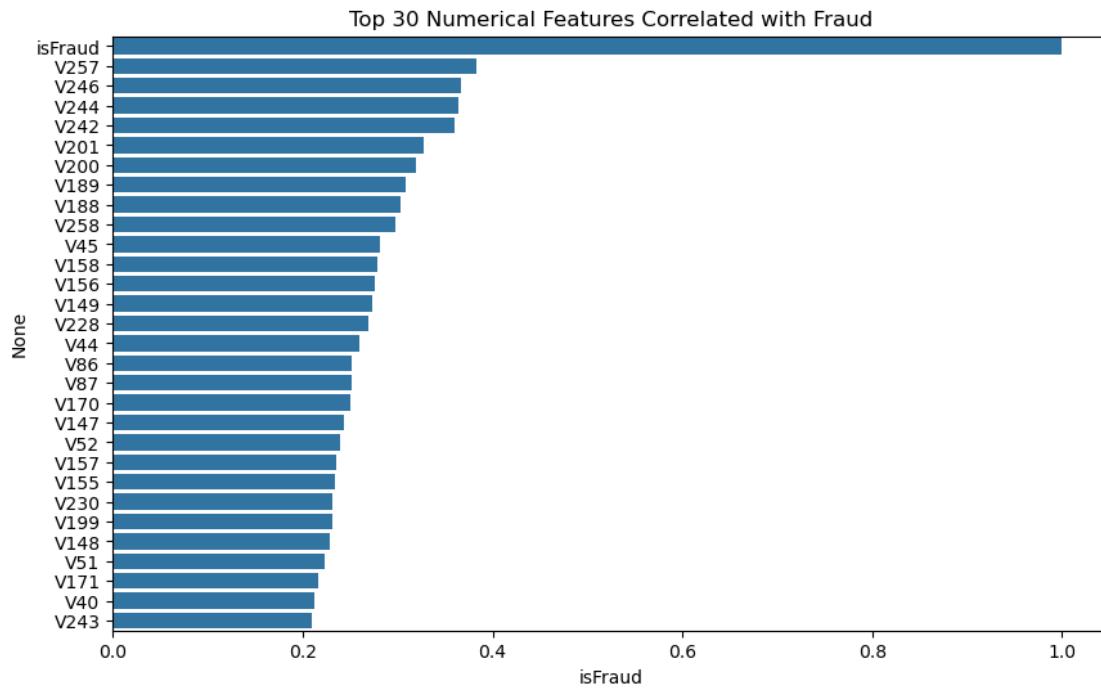


We also observe which products are being bought, seeing that product W has the largest presence in the dataset. We also observe most transactions are with `mastercard` or `visa` and most are `debit`.

```
[42]: numeric_cols = train_df.select_dtypes(include=np.number).columns
corr = train_df[numeric_cols].corr()['isFraud'].sort_values(ascending=False)
corr.head(10)
```

```
[42]: isFraud      1.000000
V257        0.383060
V246        0.366878
V244        0.364129
V242        0.360590
V201        0.328005
V200        0.318783
V189        0.308219
V188        0.303582
V258        0.297151
Name: isFraud, dtype: float64
```

```
[48]: plt.figure(figsize=(10,6))
sns.barplot(x=corr.head(30), y=corr.head(30).index)
plt.title('Top 30 Numerical Features Correlated with Fraud')
plt.show()
```



We observe that there are multiple V columns that are highly correlated with isFraud